

Navigation

January 16, 2019

1 Navigation

You are welcome to use this coding environment to train your agent for the project. Follow the instructions below to get started!

1.0.1 1. Start the Environment

Run the next code cell to install a few packages. This line will take a few minutes to run!

```
In [5]: !pip -q install ./python
```

The environment is already saved in the Workspace and can be accessed at the file path provided below. Please run the next code cell without making any changes.

```
In [6]: from unityagents import UnityEnvironment
import numpy as np

# please do not modify the line below
env = UnityEnvironment(file_name="/data/Banana_Linux_NoVis/Banana.x86_64")
```

```
INFO:unityagents:
```

```
'Academy' started successfully!
```

```
Unity Academy name: Academy
```

```
    Number of Brains: 1
```

```
    Number of External Brains : 1
```

```
    Lesson number : 0
```

```
    Reset Parameters :
```

```
Unity brain name: BananaBrain
```

```
    Number of Visual Observations (per agent): 0
```

```
    Vector Observation space type: continuous
```

```
    Vector Observation space size (per agent): 37
```

```
    Number of stacked Vector Observation: 1
```

```
    Vector Action space type: discrete
```

```
    Vector Action space size (per agent): 4
```

```
    Vector Action descriptions: , , ,
```

Environments contain *brains* which are responsible for deciding the actions of their associated agents. Here we check for the first brain available, and set it as the default brain we will be controlling from Python.

```
In [3]: # get the default brain
        brain_name = env.brain_names[0]
        brain = env.brains[brain_name]
```

1.0.2 2. Examine the State and Action Spaces

Run the code cell below to print some information about the environment.

```
In [4]: # reset the environment
        env_info = env.reset(train_mode=True)[brain_name]

        # number of agents in the environment
        print('Number of agents:', len(env_info.agents))

        # number of actions
        action_size = brain.vector_action_space_size
        print('Number of actions:', action_size)

        # examine the state space
        state = env_info.vector_observations[0]
        print('States look like:', state)
        state_size = len(state)
        print('States have length:', state_size)
```

```
Number of agents: 1
Number of actions: 4
States look like: [ 1.          0.          0.          0.          0.84408134  0.          0.
 1.          0.          0.0748472  0.          1.          0.          0.
 0.25755     1.          0.          0.          0.          0.74177343
 0.          1.          0.          0.          0.25854847  0.          0.
 1.          0.          0.09355672  0.          1.          0.          0.
 0.31969345  0.          0.          ]
States have length: 37
```

1.0.3 3. Take Random Actions in the Environment

In the next code cell, you will learn how to use the Python API to control the agent and receive feedback from the environment.

Note that in this coding environment, you will not be able to watch the agent while it is training, and you should set `train_mode=True` to restart the environment.

```
In [5]: env_info = env.reset(train_mode=True)[brain_name] # reset the environment
        state = env_info.vector_observations[0]           # get the current state
        score = 0                                         # initialize the score
```

```

while True:
    action = np.random.randint(action_size)           # select an action
    env_info = env.step(action)[brain_name]           # send the action to the environment
    next_state = env_info.vector_observations[0]        # get the next state
    reward = env_info.rewards[0]                      # get the reward
    done = env_info.local_done[0]                     # see if episode has finished
    score += reward                                    # update the score
    state = next_state                                # roll over the state to next time st
    if done:                                           # exit loop if episode finished
        break

print("Score: {}".format(score))

```

Score: 0.0

When finished, you can close the environment.

```
In [6]: env.close()
```

1.0.4 4. It's Your Turn!

Now it's your turn to train your own agent to solve the environment! A few **important notes**: - When training the environment, set `train_mode=True`, so that the line for resetting the environment looks like the following:

```
env_info = env.reset(train_mode=True)[brain_name]
```

- To structure your work, you're welcome to work directly in this Jupyter notebook, or you might like to start over with a new file! You can see the list of files in the workspace by clicking on *Jupyter* in the top left corner of the notebook.
- In this coding environment, you will not be able to watch the agent while it is training. However, *after training the agent*, you can download the saved model weights to watch the agent on your own machine!

```

In [7]: import torch
        import numpy as np
        from collections import deque
        import matplotlib.pyplot as plt
        %matplotlib inline
        #from unityagents import UnityEnvironment

        #
        from dqn_agent import Agent

```

```

In [31]: def dqn(env, n_episodes=2000, max_t=1000, eps_start=1.0, eps_end=0.01, eps_decay=0.995)
        """Deep Q-Learning.

        Params

```

```

=====
    env (Unity Environment): environment
    n_episodes (int): maximum number of training episodes
    max_t (int): maximum number of timesteps per episode
    eps_start (float): starting value of epsilon, for epsilon-greedy action selection
    eps_end (float): minimum value of epsilon
    eps_decay (float): multiplicative factor (per episode) for decreasing epsilon
"""
scores = []                                # list containing scores from each episode
scores_window = deque(maxlen=100)          # last 100 scores
eps = eps_start                             # initialize epsilon
solved = False
for i_episode in range(1, n_episodes+1):
    # reset env
    env_info = env.reset(train_mode=True)[brain_name]
    state = env_info.vector_observations[0]
    score = 0
    for t in range(max_t):
        action = agent.act(state, eps)
        env_info = env.step(action)[brain_name]
        next_state = env_info.vector_observations[0]
        reward = env_info.rewards[0]
        done = env_info.local_done[0]

        agent.step(state, action, reward, next_state, done)
        state = next_state
        score += reward
        if done:
            break
    scores_window.append(score)              # save most recent score
    scores.append(score)                    # save most recent score
    eps = max(eps_end, eps_decay*eps)      # decrease epsilon

    # Log progress
    print('\rEpisode {} \tAverage Score: {:.2f}'.format(i_episode, np.mean(scores_window)))
    if i_episode % 100 == 0:
        print('\rEpisode {} \tAverage Score: {:.2f}'.format(i_episode, np.mean(scores_window)))
    if np.mean(scores_window) >= 13.0 and not solved:
        solved = True
        print('\nEnvironment solved in {:d} episodes! \tAverage Score: {:.2f}'.format(i_episode, np.mean(scores_window)))
        torch.save(agent.qnetwork_local.state_dict(), 'checkpoint.pth')
return scores

```

```

In [32]: brain_name = env.brain_names[0]
        brain = env.brains[brain_name]
        env_info = env.reset(train_mode=True)[brain_name]
        action_size = brain.vector_action_space_size

```


```
state = env_info.vector_observations[0]
state_size = len(state)
```

```
In [33]: agent = Agent(state_size=state_size, action_size=action_size, seed=0)
```

```
In [34]: scores = dqn(env)
```

```
Episode 100      Average Score: 0.71
Episode 200      Average Score: 4.81
Episode 300      Average Score: 8.07
Episode 400      Average Score: 10.72
Episode 463      Average Score: 13.01
Environment solved in 363 episodes!      Average Score: 13.01
Episode 500      Average Score: 13.57
Episode 600      Average Score: 14.41
Episode 700      Average Score: 14.92
Episode 800      Average Score: 15.86
Episode 900      Average Score: 14.99
Episode 1000     Average Score: 14.53
Episode 1100     Average Score: 15.57
Episode 1200     Average Score: 15.32
Episode 1300     Average Score: 15.82
Episode 1400     Average Score: 15.31
Episode 1500     Average Score: 15.37
Episode 1600     Average Score: 15.95
Episode 1700     Average Score: 15.87
Episode 1800     Average Score: 16.14
Episode 1900     Average Score: 15.73
Episode 2000     Average Score: 15.14
```

```
In [35]: fig = plt.figure()
         ax = fig.add_subplot(111)
         plt.plot(np.arange(len(scores)), scores)
         plt.xlabel('Episodes')
         plt.ylabel('Score')
         plt.show()
```



In []: